

Week 5 assignment part 2

◆ 1: Problem Scope (5 points)

Problem

Hospitals face both financial and clinical challenges due to patients being readmitted within 30 days of discharge. Reducing unnecessary readmissions is a key priority for improving healthcare quality and efficiency.

Objective

Develop an AI system that predicts the likelihood of a patient being readmitted within 30 days of discharge, so healthcare staff can take preventive measures.

Stakeholders

- **Hospital Management** – reduce costs, improve quality scores.
- **Clinicians (Doctors & Nurses)** – use risk predictions to guide follow-up care.
- **Patients** – receive proactive and personalized care.
- **Data & IT Teams** – support data integration, privacy, and deployment.

◆ 2: Data Strategy (10 points)

a) Proposed Data Sources

- **Electronic Health Records (EHR)**: admission/discharge dates, diagnoses (ICD-10), procedures, lab results.
- **Patient Demographics**: age, gender, socioeconomic status.
- **Admission History**: prior readmissions, chronic conditions.
- **Medications**: prescribed drugs at discharge.

b) Ethical Concerns

1. **Patient Privacy**: Medical data is sensitive and must be protected (HIPAA compliance).

2. **Algorithmic Bias:** Data may underrepresent certain populations (e.g., rural or minority patients), leading to unfair predictions.

◆ 3: Model Development (10 points)

a) Model Selection and Justification

Random Forest Classifier

- Handles mixed data types (numerical & categorical).
- Reduces overfitting through ensemble learning.
- Provides feature importance for interpretability.

◆ 4: Deployment (10 points)

a) Integration into Hospital Systems

1. **Model Packaging:** Save model using joblib or pickle.
2. **API Deployment:** Use **Flask** or **FastAPI** to expose predictions via REST API.
3. **EHR Integration:** Connect to hospital systems like Epic or Cerner to run predictions at discharge.
4. **Clinician Dashboard:** Show readmission risk score and top contributing features.
5. **Monitoring:** Log predictions and flag drift or unexpected results.

b) Regulatory Compliance

- **Data Encryption:** Secure data at rest and in transit (TLS/HTTPS).
- **Access Control:** Ensure only authorized staff can access predictions.
- **Audit Logging:** Record access to model inputs/outputs.
- **De-identification:** Remove or mask sensitive data used for training.
- **HIPAA Compliance:** Work with legal teams to align with healthcare regulations

◆ 5: Optimization (5 points)

🔧 Method to Address Overfitting:

Cross-Validation + Regularization

- Use **K-Fold Cross-Validation** to ensure the model generalizes across different subsets of data.
- In tree-based models like Random Forest:
 - Limit max_depth
 - Increase min_samples_split
 - Tune n_estimators